Logistic Regression Project

In this project we will be working with a fake advertising data set, indicating whether or not a particular internet user clicked on an Advertisement. We will try to create a model that will predict whether or not they will click on an ad based off the features of that user.

This data set contains the following features:

- 'Daily Time Spent on Site': consumer time on site in minutes
- 'Age': cutomer age in years
- 'Area Income': Avg. Income of geographical area of consumer
- 'Daily Internet Usage': Avg. minutes a day consumer is on the internet
- 'Ad Topic Line': Headline of the advertisement
- 'City': City of consumer
- 'Male': Whether or not consumer was male
- 'Country': Country of consumer
- 'Timestamp': Time at which consumer clicked on Ad or closed window
- 'Clicked on Ad': 0 or 1 indicated clicking on Ad

Import Libraries

Import a few libraries you think you'll need (Or just import them as you go along!)

In [2]:
```
import pandas as pd
import numpy as np
```

In [3]:
```
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Get the Data

Read in the advertising.csv file and set it to a data frame called ad_data.

In [4]:
```
ad_data = pd.read_csv('advertising.csv')
```

Check the head of ad_data

In [5]:
```
ad_data.head()
```

Out[5]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Ad Topic Line	City	Male	Country	Timestamp	Clicked on Ad
0	68.95	35	61833.90	256.09	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	2016-03-27 00:53:11	0
1	80.23	31	68441.85	193.77	Monitored national standardization	West Jodi	1	Nauru	2016-04-04 01:39:02	0
2	69.47	26	59785.94	236.50	Organic bottom-line service-desk	Davidton	0	San Marino	2016-03-13 20:35:42	0
3	74.15	29	54806.18	245.89	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	2016-01-10 02:31:19	0
4	68.37	35	73889.99	225.58	Robust logistical utilization	South Manuel	0	Iceland	2016-06-03 03:36:18	0

Use info and describe() on ad_data

In [6]:
```
ad_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
 #   Column                    Non-Null Count  Dtype  
---  ------                    --------------  -----  
 0   Daily Time Spent on Site  1000 non-null   float64
 1   Age                       1000 non-null   int64  
 2   Area Income               1000 non-null   float64
 3   Daily Internet Usage      1000 non-null   float64
 4   Ad Topic Line             1000 non-null   object 
 5   City                      1000 non-null   object 
 6   Male                      1000 non-null   int64  
 7   Country                   1000 non-null   object 
 8   Timestamp                 1000 non-null   object 
 9   Clicked on Ad             1000 non-null   int64  
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

In [8]:
```
ad_data.describe()
```

Out[8]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Male	Clicked on Ad
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.00000
mean	65.000200	36.009000	55000.000080	180.000100	0.481000	0.50000
std	15.853615	8.785562	13414.634022	43.902339	0.499889	0.50025
min	32.600000	19.000000	13996.500000	104.780000	0.000000	0.00000
25%	51.360000	29.000000	47031.802500	138.830000	0.000000	0.00000
50%	68.215000	35.000000	57012.300000	183.130000	0.000000	0.50000
75%	78.547500	42.000000	65470.635000	218.792500	1.000000	1.00000
max	91.430000	61.000000	79484.800000	269.960000	1.000000	1.00000

Exploratory Data Analysis

Let's use seaborn to explore the data!

Try recreating the plots shown below!

Create a histogram of the Age

In [14]:
```
sns.histplot(x = 'Age',data = ad_data,bins=30)
```

Out[14]:
```
<AxesSubplot:xlabel='Age', ylabel='Count'>
```

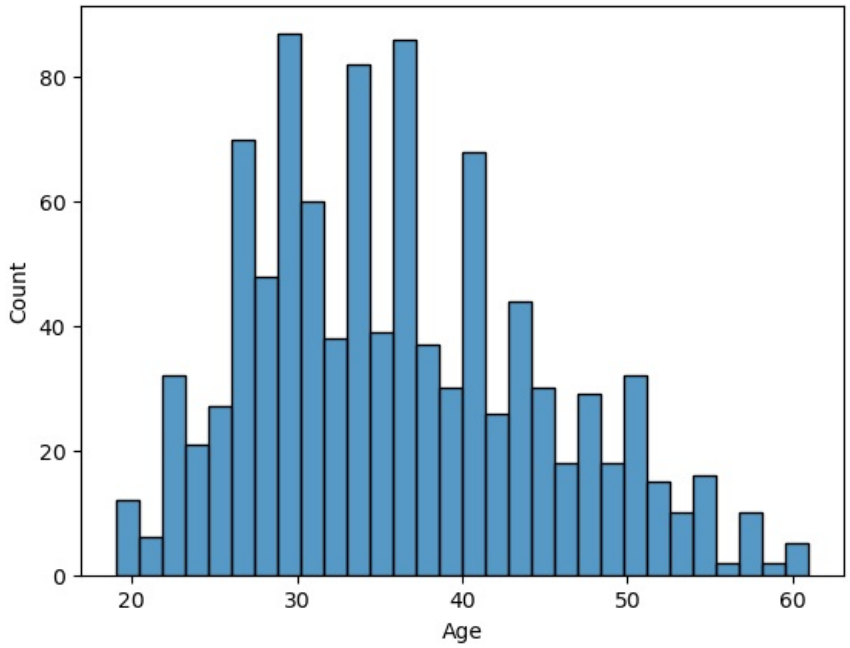

Create a jointplot showing Area Income versus Age.

In [13]:
```
sns.jointplot(x= 'Age',y='Area Income',data= ad_data,)
```

Out[13]:
```
<seaborn.axisgrid.JointGrid at 0x1dc6f43e580>
```

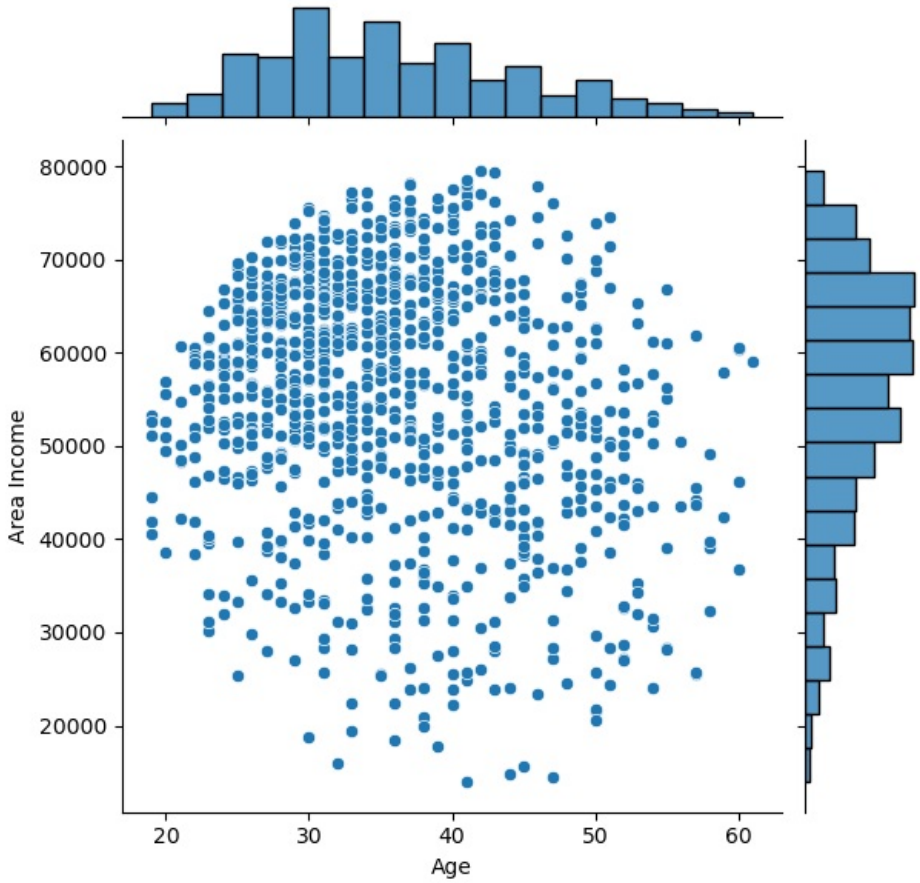

Create a jointplot showing the kde distributions of Daily Time spent on site vs. Age.

In [20]:
```
sns.jointplot(x='Age',y='Daily Time Spent on Site',data = ad_data,kind='kde',color='red')

#Like Area Income, there is a high concentration of points at high levels between ages 30 and 40.
# In that age gap, there is more time spent daily on the site and the area income is higher
```

Out[20]:
```
<seaborn.axisgrid.JointGrid at 0x1dc70a99bb0>
```

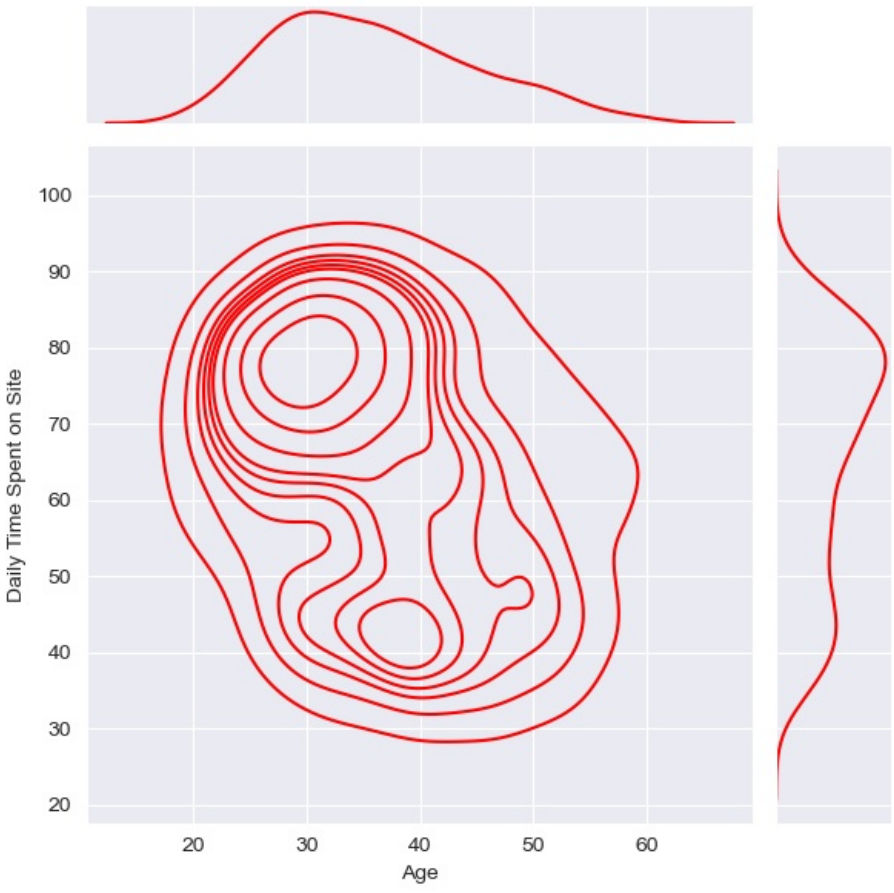

Create a jointplot of 'Daily Time Spent on Site' vs. 'Daily Internet Usage'

In [21]:
```
sns.jointplot(x='Daily Time Spent on Site', y='Daily Internet Usage',data=ad_data)
```

Out[21]:
```
<seaborn.axisgrid.JointGrid at 0x1dc72aae190>
```

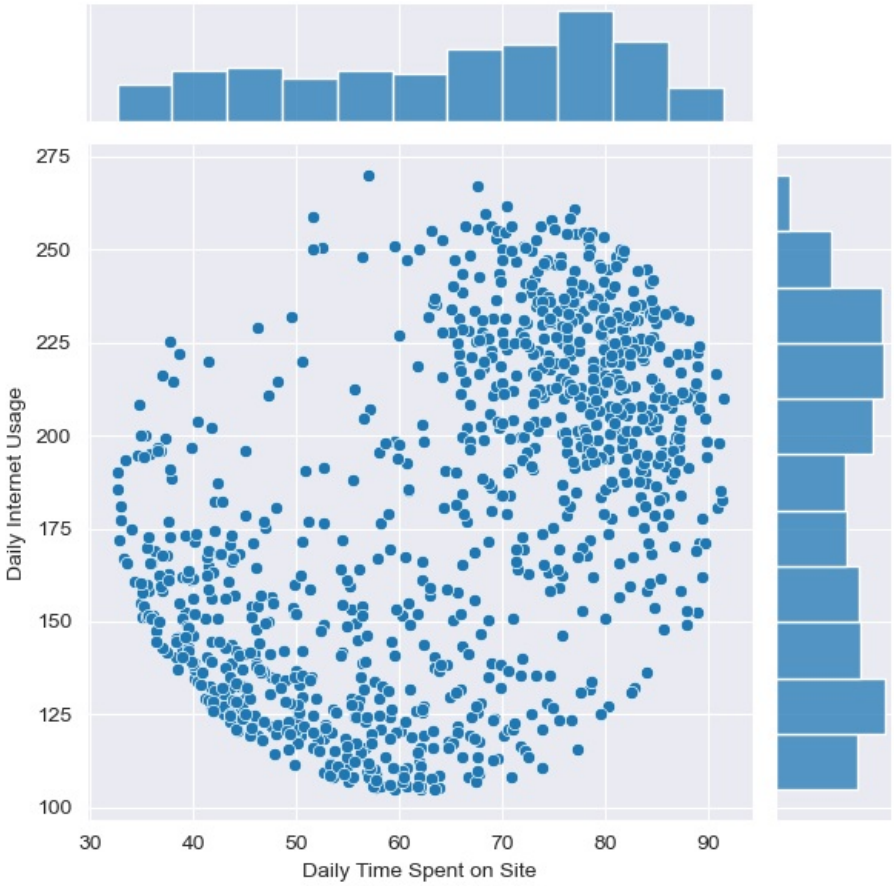

Finally, create a pairplot with the hue defined by the 'Clicked on Ad' column feature.

In [22]:
```
sns.pairplot(ad_data,hue='Clicked on Ad')
```

Out[22]:
```
<seaborn.axisgrid.PairGrid at 0x1dc730395b0>
```

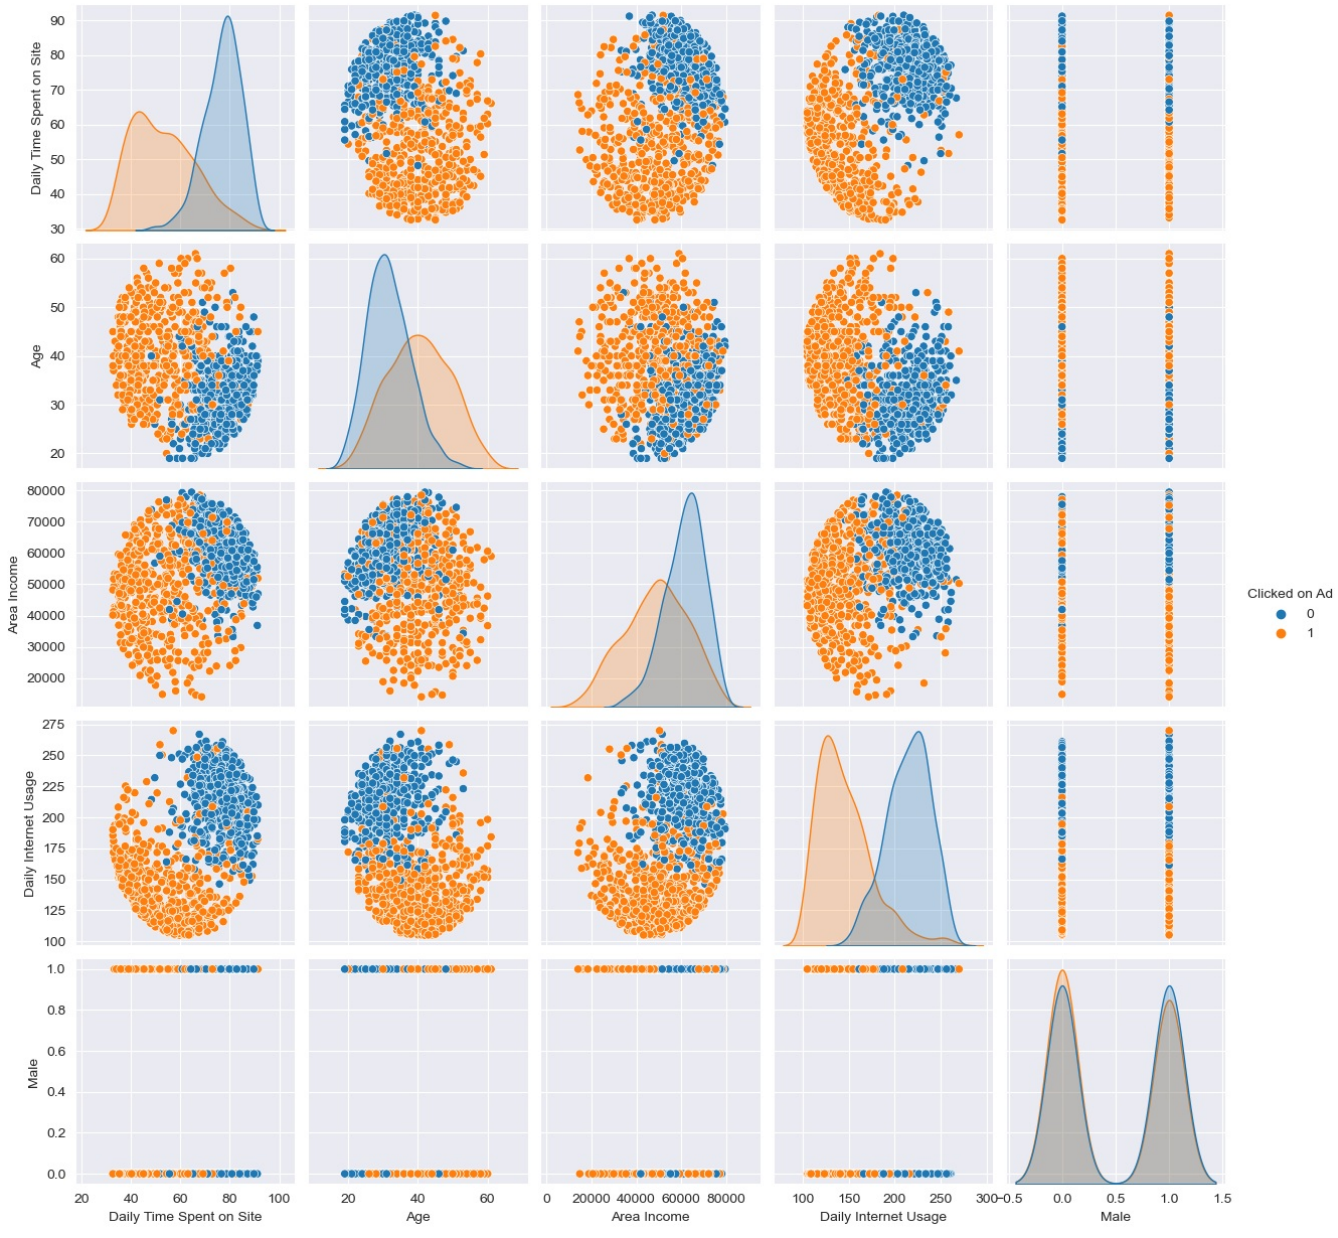
Daily Time Spent on Site
Age
Area Income
Daily Internet Usage
Male
Clicked on Ad
0
1

Logistic Regression

Now it's time to do a train test split, and train our model!

You'll have the freedom here to choose columns that you want to train on!

Split the data into training set and testing set using train_test_split

In [24]:
```
from sklearn.model_selection import train_test_split
```

In [27]:
```
#Let's define our X and y
# The columns that are not picked just will not fit in our logsitic regression model

X = ad_data[['Daily Time Spent on Site','Age','Area Income','Daily Internet Usage','Male']]

# This is our target variable
y= ad_data['Clicked on Ad']


#Taken from shift+tab - train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

Train and fit a logistic regression model on the training set.

In [29]:
```
from sklearn.linear_model import LogisticRegression
```

In [30]:
```
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
```

Out[30]:
```
LogisticRegression()
```

Predictions and Evaluations

Now predict values for the testing data.

In [31]:
```
predictions = logmodel.predict(X_test)
```

Create a classification report for the model.

In [33]:
```
from sklearn.metrics import classification_report,confusion_matrix
print(classification_report(y_test,predictions))
print(confusion_matrix(y_test,predictions))
```

```
              precision    recall  f1-score   support

           0       0.91      0.95      0.93       157
           1       0.94      0.90      0.92       143

    accuracy                           0.93       300
   macro avg       0.93      0.93      0.93       300
weighted avg       0.93      0.93      0.93       300

[[149   8]
 [ 14 129]]
```

Great Job!